

# Project Appendix

Zhichao Yang, 20661179 & Liwen Dai, 20552153

## 1 Setups

Setup HDFS on pseudo distributed mode, with 2 replicas for each file, block size 64 MB:

hadoop-2.7.2/etc/hadoop/core-site.xml:

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

hadoop-2.7.2/etc/hadoop/hdfs-site.xml:

```
<configuration>
<property>
  <name>dfs.replication</name>
  <value>2</value>
</property>
<property>
  <name>dfs.heartbeat.interval</name>
  <value>500</value>
</property>
<property>
  <name>dfs.namenode.name.dir</name>
  <value>file:/<path>/namenode</value>
</property>
<property>
  <name>dfs.datanode.data.dir</name>
  <value>file:/<path>/datanode</value>
</property>
<property>
  <name>dfs.block.size</name>
  <value>67108864</value>
  <description>Block size</description>
</property>
</configuration>
```

Also need to configure ssh key to login localhost without password:

[https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-common/SingleCluster.html#Pseudo-Distributed\\_Operation](https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-common/SingleCluster.html#Pseudo-Distributed_Operation)

Format the NameNode before launch:

```
$ hdfs namenode -format
```

Web interface for HDFS state <http://localhost:50070/>

To start HDFS,

```
$ cd hadoop-2.7.2/sbin
$ ./start-dfs.sh
$ jps
```

Now should see DataNode, NameNode, SecondaryNameNode and Jps.

To run additional data nodes, use run-additionalDN.sh, modified from:

<https://bigdata.wordpress.com/2010/05/27/hadoop-cookbook-4-how-to-run-multiple-data-nodes-on-on-e-machine/>

```
#!/bin/sh
# This is used for starting multiple datanodes on the same machine.
# run it from hadoop-dir/ just like 'bin/hadoop'

#Usage: run-additionalDN.sh [start|stop] dnnumber
#e.g. run-datanode.sh start 2

DN_DIR_PREFIX="/home/<username>/mydata/hadoopTmp/dn"

if [ -z $DN_DIR_PREFIX ]; then
echo $0: DN_DIR_PREFIX is not set. set it to something like "/hadoopTmp/dn"
exit 1
fi

run_datanode () {
DN=$2
export HADOOP_LOG_DIR=$DN_DIR_PREFIX$DN/logs
export HADOOP_PID_DIR=$HADOOP_LOG_DIR
DN_CONF_OPTS="\
-Dhadoop.tmp.dir=$DN_DIR_PREFIX$DN"
sbin/hadoop-daemon.sh --script bin/hdfs $1 datanode $DN_CONF_OPTS
}

cmd=$1
shift;

for i in $*
do
run_datanode $cmd $i
done
```

Before use modify the DN\_DIR\_PREFIX in run-additionalDN.sh to your path.

To run additional DataNodes, need multiple copies of hadoop directories.

```
$ cp -r hadoop-2.7.2/ node1
$ cp run-additionalDN.sh node1/
```

Modify node1/etc/hadoop/hdfs-site.xml to use different directory and port numbers:

```
<configuration>
<property>
  <name>dfs.datanode.data.dir</name>
  <value>file:/home/singulo/mydata/hdfs/datanode1</value>
</property>
<property>
  <name>dfs.heartbeat.interval</name>
  <value>500</value>
</property>
<property>
  <name>dfs.datanode.address</name>
  <value>localhost:50110</value>
</property>
  <property>
    <name>dfs.datanode.ipc.address</name>
    <value>localhost:50120</value>
  </property>
  <property>
    <name>dfs.datanode.http.address</name>
    <value>localhost:50130</value>
  </property>
</configuration>
```

```
$ cd node1
$ ./run-additionalDN.sh start 1
```

Now run \$ jps again , should see an extra DataNode.

To run more data nodes, repeat the steps above, and don't forget to use new name/directory/ports.

## 2. Tool Used:

We are using the unix tool **strace** to attain system calls on namenode and datanode. The tool strace provides various utilities that are very useful in getting time-based and count-based system call information to help us understand the file system access pattern in HDFS. The -p option in strace can attach the strace to the node process, and -f option can help trace all the child processes created by a node process. -c option can count all the system call upon the end of strace. -e option can help us filter the results we want to see.

Specifically, we have the following script to help automating our tests:

```
#!/bin/bash

rm -rf trace-hdfs-proc-out
mkdir trace-hdfs-proc-out

PASSWORD="your pass word"

jps | while read node ; do
    pid=$(echo $node | cut -d' ' -f1)
    nodeName=$(echo $node | cut -d' ' -f2)
    echo $PASSWORD | sudo -S strace $1 -t -e trace=file,desc -f -o
    trace-hdfs-proc-out/${pid}.${nodeName}.txt -p $pid &
done

sleep 5
read -n 1 -p "Press any key to stop:" key

echo $PASSWORD | sudo -S kill $(ps -aux | grep strace | tr -s ' ' | cut -d' ' -f2 )
```

### 3. Results:

#### 3.1 File system access pattern when HDFS is idle:

Search system-wide file descriptor:

```
$ lsof | grep ...
```

Or

```
$ readlink /proc/<pid>/fd/<fd number>
```

##### 3.1.1 Each of 3 DataNodes:

Repeatedly call `epoll_wait()`, on several file descriptors, where `readlink` returns `anon_inode:[eventpoll]` (what is this?)

Loop, heartbeat to NameNode:

- 3 calls to `statfs()` on the root of node
- Try `read()` from TCP = -1 EAGAIN
- `EPOLL_CTL_ADD`
- `write()` to TCP fd, 398 bytes, (talk to NameNode)
- `EPOLL_CTL_DEL`
- Read 40 bytes from TCP
- `Epoll_wait()`

Every loop lasts for about 3 seconds.

Notices that there's no writes to the disc.

The default heartbeat time is 3 seconds, change it to 500 in the configuration file to remove the "heartbeat noise" in the result. In every `hdfs-site.xml` file, add:

```
<property>
  <name>dfs.heartbeat.interval</name>
  <value>500</value>
  <description>Heartbeat interval</description>
</property>
```

### 3.1.2 NameNode:

Loop:

statfs() on the node root

Epoll\_wait

From each DataNode, read 4 + 394 = 398 bytes from a TCP socket, the reply 40 bytes

Each loop lasts about 5 seconds.

Besides, Fcntl F\_GETFL, F\_GETFL appeared 3 times

### 3.1.3 Secondary NameNode:

Log of trace is only 1-page long, mostly epoll\_wait(), with few small reads/writes, epoll\_ctl, fcntl, dup2, mmap(size=12288)

Understanding what HDFS is doing at idle state helps to remove “noise” at real workload analysis.

## 3.2 File system access pattern when appending short text to a file:

We keep appending a text file of 5.1 MB, Shakespeare.txt, to the HDFS.

```
#!/bin/bash
for i in $(seq 1 50) ; do hdfs dfs -appendToFile Shakespeare.txt Shakespeare-copy.txt ; done
```

### 3.2.1 DataNode:

Before writing the actual content, a datanode checks stats of, and reads from many java class, jar and xml files, and it keeps calling lseek() and read() on hadoop-hdfs-2.7.2.jar (fd = 107), where most reads are 30 bytes

```
25644 19:22:32
stat("/home/singulo/hadoop-2.7.2/etc/hadoop/org/apache/hadoop/hdfs/protocol/proto/HdfsProtos$BlockProto.class", 0x7fad8be3f200) = -1 ENOENT (No such file or directory)
25644 19:22:32
stat("/home/singulo/hadoop-2.7.2/share/hadoop/hdfs/org/apache/hadoop/hdfs/protocol/proto/HdfsProtos$BlockProto.class", 0x7fad8be3f200) = -1 ENOENT (No such file or directory)
25644 19:22:32 lseek(107, 469827, SEEK_SET) = 469827
25644 19:22:32 read(107, "PK\3\4\24\0\10\0\10\0Z\1:H\0\0\0\0\0\0\0\0\0\0A\0\0\0", 30) = 30
25644 19:22:32 lseek(107, 469922, SEEK_SET) = 469922
25644 19:22:32 read(107, "\275Z\vx\24\327u\376\357jwGZFBO@,\n\2\213\225@\17\236F26BB,"..., 4463) = 4463
25644 19:22:32 mmap(0x7fad8a3ee000, 262144, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x7fad8a3ee000
```

During this process, the block files blk\_xxx and blk\_xxx.meta are created under BP-xxx/current/rbw/ directory, which represents the “Replica Being Written” state.

```
2893 19:22:32
stat("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/curr
ent/rbw/blk_1073741827", 0x7fad8a62d200) = -1 ENOENT (No such file or directory)
2893 19:22:32
open("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/cur
rent/rbw/blk_1073741827", O_RDWR|O_CREAT|O_EXCL, 0666) = 324
2893 19:22:32 fstat(324, {st_mode=S_IFREG|0664, st_size=0, ...}) = 0
2893 19:22:32 close(324)
...
2893 19:22:32
open("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/cur
rent/rbw/blk_1073741827_1042.meta", O_RDWR|O_CREAT, 0666) = 324
2893 19:22:32 fstat(324, {st_mode=S_IFREG|0664, st_size=0, ...}) = 0
2893 19:22:32
open("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/cur
rent/rbw/blk_1073741827", O_RDWR|O_CREAT, 0666) = 325
2893 19:22:32 fstat(325, {st_mode=S_IFREG|0664, st_size=0, ...}) = 0
```

64 KBytes at a time, the datanode reads the data from a socket (fd = 316), and writes to another socket (fd=326), to another replica. Then it appends the data to the blk file, and appends 504 bytes to the blk.meta file.

```
2893 19:22:32 read(316,
"\0\0\375\374\0\31\t\0\340\7\0\0\0\0\0\21\10\0\0\0\0\0\0\30\0%\0\374\0\0\305"... , 512) = 512
2893 19:22:32 read(316, "5"Y&9\222$\352\3138\25\373\277\335\2\251b\357\257,b\33\367 of
fores"... , 64535) = 64535
2893 19:22:32 write(326,
"\0\0\375\374\0\31\t\0\340\7\0\0\0\0\0\21\10\0\0\0\0\0\0\30\0%\0\374\0\0\305"... , 65047) = 65047
2893 19:22:32 write(325, " of foresters\n\n JAQUES. Which i"... , 64512) = 64512
2893 19:22:32 write(324,
"\305\4W\375\0374\327\10k\275u\1.\236\326^\326\30\3\377\310\271;\257\4\1\332\260\331\201\1#"
..., 504) = 504
...
```

When the actual data have been written, the datanode rename the rbw/blk and rbw/blk\_.meta files to the finalized/ directory. Then it writes to its own log file (357 bytes), and acknowledges back through the socket.

```
2896 19:22:32
stat("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/curr
ent/finalized/subdir0/subdir0", {st_mode=S_IFDIR|0775, st_size=4096, ...}) = 0
2896 19:22:32
rename("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/
current/rbw/blk_1073741827_1042.meta",
"/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/current/fi
nalized/subdir0/subdir0/blk_1073741827_1042.meta") = 0
2896 19:22:32
rename("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/
current/rbw/blk_1073741827",
"/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/current/fi
nalized/subdir0/subdir0/blk_1073741827") = 0
2896 19:22:32
```

```

stat("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/curr
ent/finalized/subdir0/subdir0/blk_1073741827_1042.meta", {st_mode=S_IFREG|0664,
st_size=41635, ...}) = 0
...
2896 19:22:32 write(191, "2017-03-29 19:22:32,925 INFO org"..., 357 <unfinished ...>
...
2896 19:22:32 <... write resumed> ) = 357
2896 19:22:32 write(316, "\20\10\246\1\20\0\20\0\30\277\324\252\1\\"\2\0\0", 17 <unfinished ...>
...
2896 19:22:32 <... write resumed> ) = 17

```

After that, the datanode calls close() to more than 300 file descriptors:

```

...
2986 19:22:34 close(319)          = 0
2986 19:22:34 close(320)          = 0
2986 19:22:34 close(321)          = 0
2986 19:22:34 close(322)          = 0
2986 19:22:34 close(323)          = 0
2986 19:22:34 close(327)          = 0
2986 19:22:34 close(328)          = 0
2986 19:22:34 close(329)          = 0

```

When the next append command comes, finalized/blk\_xxx(.meta) are renamed back to rbw/ directory, and the pattern above repeats.

```

2985 19:22:34 statfs("/home/singulo/mydata/hdfs/datanode", {f_type="EXT2_SUPER_MAGIC",
f_bsize=4096, f_blocks=27892164, f_bfree=17804223, f_bavail=16381619, f_files=7094272,
f_ffree=6501723, f_fsid={-1149700033, -846037428}, f_namelen=255, f_frsize=4096,
f_flags=4128}) = 0
2985 19:22:34 statfs("/home/singulo/mydata/hdfs/datanode", {f_type="EXT2_SUPER_MAGIC",
f_bsize=4096, f_blocks=27892164, f_bfree=17804223, f_bavail=16381619, f_files=7094272,
f_ffree=6501723, f_fsid={-1149700033, -846037428}, f_namelen=255, f_frsize=4096,
f_flags=4128}) = 0
2985 19:22:34
rename("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/
current/finalized/subdir0/subdir0/blk_1073741827_1042.meta",
"/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/current/r
bw/blk_1073741827_1043.meta") = 0
2985 19:22:34
rename("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/
current/finalized/subdir0/subdir0/blk_1073741827",
"/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/current/r
bw/blk_1073741827") = 0
2985 19:22:34
open("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/cur
rent/rbw/blk_1073741827_1043.meta", O_RDWR|O_CREAT, 0666) = 324
2985 19:22:34 fstat(324, {st_mode=S_IFREG|0664, st_size=41635, ...}) = 0
....

```

When the data file size exceeds 64 MB, the block size we set, the new rbw/blk files are created accordingly.

Call Summary from a DataNode:

% time	seconds	usecs/call	calls	errors	syscall
99.52	0.496050	432	1148		epoll_wait
0.43	0.002119	1	4043	28	write
0.05	0.000258	0	4159	445	read
0.00	0.000018	0	666		mmap
0.00	0.000012	0	968		epoll_ctl
0.00	0.000000	0	1107	792	open
0.00	0.000000	0	10355		close
0.00	0.000000	0	203	88	stat
0.00	0.000000	0	495		fstat
0.00	0.000000	0	220		lstat
0.00	0.000000	0	98		lseek
0.00	0.000000	0	144		ioctl
0.00	0.000000	0	199	199	access
0.00	0.000000	0	147		pipe
0.00	0.000000	0	141		dup2
0.00	0.000000	0	232	203	execve
0.00	0.000000	0	191		fcntl
0.00	0.000000	0	72		getdents
0.00	0.000000	0	58		rename
0.00	0.000000	0	95	56	statfs
0.00	0.000000	0	2		epoll_create
0.00	0.000000	0	7		openat
0.00	0.000000	0	31		newfstatat
100.00	0.498457		24781	1811	total

### 3.2.2 NameNode:

The NameNode writes to its own log (fd = 191), with about 200 bytes each time, and talks through socket (fd = 223).

```

25465 19:23:12 write(223,
"\0\0\0005\32\10\3\20\0\30\t:\20\272\367\331\34\334\236LP\241N\326\36\26\4\202i@\0\31"... , 57)
= 57
25452 19:23:12 <... epoll_wait resumed> [{EPOLLIN, {u32=223, u64=8340593504036061407}}],
8192, -1) = 1
25452 19:23:12 read(223, "\0\0\1\364", 4) = 4
25452 19:23:12 read(223,
"\32\10\2\20\0\30\10"\20\272\367\331\34\334\236LP\241N\326\36\26\4\202i(\0B\n\16up"... , 500) =
500
25452 19:23:12 epoll_wait(215, <unfinished ...>
25473 19:23:12 write(191, "2017-03-29 19:23:12,826 INFO org"... , 240) = 240
25473 19:23:12 write(191, "2017-03-29 19:23:12,826 INFO Blo"... , 161) = 161
25473 19:23:12 write(191, "2017-03-29 19:23:12,826 INFO Blo"... , 161) = 161

```



Similar to DataNode, the NameNode also performs lots reads and stats checks on various code files (.class, libnio.so, etc...). However, no rename() is observed, and the actual data doesn't go through the NameNode.

The NameNode call summary:

% time	seconds	usecs/call	calls	errors	syscall
96.40	0.431255	583	740		epoll_wait
2.68	0.012000	59	204		fdatasync
0.92	0.004097	4	1047		write
0.00	0.000000	0	1245	57	read
0.00	0.000000	0	11		open
0.00	0.000000	0	67		close
0.00	0.000000	0	216		fstat
0.00	0.000000	0	226		lseek
0.00	0.000000	0	27		mmap
0.00	0.000000	0	55		dup2
0.00	0.000000	0	171		fcntl
0.00	0.000000	0	18		statfs
0.00	0.000000	0	112	55	epoll_ctl
100.00	0.447352		4139	112	total

### 3.3 File system access pattern for large writes:

We try to put out.actor-collaboration, a text file of 452 MB, to HDFS twice:

```
#!/bin/bash
hdfs dfs -put out.actor-collaboration out.actor-collaboration-1
hdfs dfs -put out.actor-collaboration out.actor-collaboration-2
```

The NameNode turns out pretty much similar to the case of short writes.

% time	seconds	usecs/call	calls	errors	syscall
86.38	0.076243	641	119		epoll_wait
13.60	0.012000	600	20		fdatasync
0.02	0.000017	1	18		fcntl
0.00	0.000000	0	165	6	read
0.00	0.000000	0	132		write
0.00	0.000000	0	5		close
0.00	0.000000	0	20		fstat
0.00	0.000000	0	20		lseek
0.00	0.000000	0	5		dup2
0.00	0.000000	0	8		statfs
0.00	0.000000	0	11	5	epoll_ctl
100.00	0.088260		523	11	total

DataNodes are also similar to the case of short writes, except that the read/write operations are much more intensive.

```
161113 23:19:44 statfs("/home/singulo/mydata/hdfs/datanode1", {f_type="EXT2_SUPER_MAGIC",
f_bsize=4096, f_blocks=27892164, f_bfree=17413536, f_bavail=15990932, f_files=7094272,
```

```

f_ffree=6496257, f_fsid={-1149700033, -846037428}, f_namelen=255, f_frsize=4096,
f_flags=4128}) = 0
16113 23:19:44 statfs("/home/singulo/mydata/hdfs/datanode1", {f_type="EXT2_SUPER_MAGIC",
f_bsize=4096, f_blocks=27892164, f_bfree=17413536, f_bavail=15990932, f_files=7094272,
f_ffree=6496257, f_fsid={-1149700033, -846037428}, f_namelen=255, f_frsize=4096,
f_flags=4128}) = 0
16113 23:19:44
stat("/home/singulo/mydata/hdfs/datanode1/current/BP-1586938737-127.0.1.1-1490573390385/cur
rent/rbw/blk_1073741839", 0x7f2d779a8200) = -1 ENOENT (No such file or directory)
16113 23:19:44
open("/home/singulo/mydata/hdfs/datanode1/current/BP-1586938737-127.0.1.1-1490573390385/c
urrent/rbw/blk_1073741839", O_RDWR|O_CREAT|O_EXCL, 0666) = 319
16113 23:19:44 fstat(319, {st_mode=S_IFREG|0664, st_size=0, ...}) = 0
16113 23:19:44 close(319) = 0
16113 23:19:44
open("/home/singulo/mydata/hdfs/datanode1/current/BP-1586938737-127.0.1.1-1490573390385/c
urrent/rbw/blk_1073741839_1103.meta", O_RDWR|O_CREAT, 0666) = 319
16113 23:19:44 fstat(319, {st_mode=S_IFREG|0664, st_size=0, ...}) = 0
16113 23:19:44
open("/home/singulo/mydata/hdfs/datanode1/current/BP-1586938737-127.0.1.1-1490573390385/c
urrent/rbw/blk_1073741839", O_RDWR|O_CREAT, 0666) = 320
...

16113 23:19:44 read(317,
"X\7\246\236\17\355\32N\226\242\266[4\vp8;FNV?1:\225\275\315\327\314L\367mc"..., 64570) =
64570
16113 23:19:44 write(326,
"\0\0\375\374\0\31\t\0\2\0\0\0\0\0\0\21\24\4\0\0\0\0\0\0\30\0%\0\374\0\0\235"..., 65047) = 65047
16113 23:19:44 write(320, "\n11419 8415 \n11419 8896 \n11419 1"..., 64512) = 64512
16113 23:19:44 write(319,
"\235lh\211\230\267\242j\200\250\311j\25\362F\232^\4\313j\357\335\272"\336\226\304C\251\\20
45"..., 504) = 504
16113 23:19:44 read(317,
"\0\0\375\374\0\31\t\0\376\0\0\0\0\0\0\21\25\4\0\0\0\0\0\0\30\0%\0\374\0\0\25"..., 512) = 512
16113 23:19:44 read(317, "h\335\3111\310\356\4J\303\0302+\24YLj\231\301\362\327\224#752
11661"..., 64535) = 64535
...

```

% time	seconds	usecs/call	calls	errors	syscall
99.73	1.598498	380	4204		epoll_wait
0.20	0.003261	0	19441	57	write
0.06	0.001041	0	17612	2007	read
0.00	0.000072	0	4135		epoll_ctl
0.00	0.000000	0	18		open
0.00	0.000000	0	30		close
0.00	0.000000	0	18	6	stat
0.00	0.000000	0	18		fstat
0.00	0.000000	0	27		mmap
0.00	0.000000	0	1		pipe
0.00	0.000000	0	9		dup2
0.00	0.000000	0	40		fcntl
0.00	0.000000	0	12		rename
0.00	0.000000	0	15		statfs

0.00	0.000000	0	1	epoll_create
-----				
100.00	1.602872		45581	2070 total

Note: `epoll_wait()` happens when read from a socket returns an error, for example:

```
15983 23:19:34 read(318, 0x7fad98221820, 512) = -1 EAGAIN (Resource temporarily unavailable)
15983 23:19:34 epoll_ctl(329, EPOLL_CTL_ADD, 318, {EPOLLIN, {u32=318, u64=4294967614}})
= 0
15983 23:19:34 epoll_wait(329, [{EPOLLIN, {u32=318, u64=4294967614}}], 8192, 60000) = 1
15983 23:19:34 epoll_ctl(329, EPOLL_CTL_DEL, 318, 0x7fad8a62d2e0) = 0
15983 23:19:34 epoll_wait(329, [], 8192, 0) = 0
15983 23:19:34 read(318, "\0\34P\252\1\n\nA\n5\n%BP-1586938737-127.0"... , 512) = 175
```

### 3. 4 File system access pattern for reads:

```
#!/bin/bash

hdfs dfs -cat out.actor-collaboration-1 > /dev/null
hdfs dfs -cat out.actor-collaboration-2 > /dev/null
```

**3.4.1 NameNode:** most are short reads/writes to the sockets. No writes to local file system are observed.

Reads `lib/guava-11.0.2.jar` etc...

% time	seconds	usecs/call	calls	errors	syscall
-----					
100.00	0.036000	1161	31		epoll_wait
0.00	0.000000	0	42	4	read
0.00	0.000000	0	12		write
0.00	0.000000	0	2		close
0.00	0.000000	0	1		mmap
0.00	0.000000	0	2		dup2
0.00	0.000000	0	12		fcntl
0.00	0.000000	0	2		statfs
0.00	0.000000	0	6	2	epoll_ctl
-----					
100.00	0.036000		110		6 total

### 3.4.2 DataNode:

% time	seconds	usecs/call	calls	errors	syscall
-----					
99.92	0.280019	50	5632		epoll_wait
0.05	0.000132	0	2801		sendfile
0.02	0.000066	0	5614		epoll_ctl
0.01	0.000022	0	2809		write
0.00	0.000008	0	367	5	read
0.00	0.000006	0	3157		fstat
0.00	0.000000	0	6		open

0.00	0.000000	0	7	close
0.00	0.000000	0	9	stat
0.00	0.000000	0	353	lseek
0.00	0.000000	0	3	mmap
0.00	0.000000	0	1	dup2
0.00	0.000000	0	9	fcntl
0.00	0.000000	0	3	statfs
0.00	0.000000	0	91	fadvise64
-----				
100.00	0.280253		20862	5 total

For each block, only 1 replica is read. First, the block and the meta files are opened:

```

31179 16:04:22 read(316, "\0\34Qt\ng\nA\n5\n%BP-1586938737-127.0."..., 512) = 120
31179 16:04:22
stat("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/curr
ent/finalized/subdir0/subdir0/blk_1073741850", {st_mode=S_IFREG|0664, st_size=67108864, ...})
= 0
31179 16:04:22
stat("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/curr
ent/finalized/subdir0/subdir0/blk_1073741850_1114.meta", {st_mode=S_IFREG|0664,
st_size=524295, ...}) = 0
31179 16:04:22
open("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/cur
rent/finalized/subdir0/subdir0/blk_1073741850_1114.meta", O_RDONLY) = 318
31179 16:04:22 fstat(318, {st_mode=S_IFREG|0664, st_size=524295, ...}) = 0
31179 16:04:22
stat("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/curr
ent/finalized/subdir0/subdir0/blk_1073741850_1114.meta", {st_mode=S_IFREG|0664,
st_size=524295, ...}) = 0
31179 16:04:22 read(318, "\0\1\2\0\0\2\0\324\35\252wt\372\325#
[\335\"y\3\305\274c>\323\242\353\t\334\24"..., 4096) = 4096
31179 16:04:22
open("/home/singulo/mydata/hdfs/datanode/current/BP-1586938737-127.0.1.1-1490573390385/cur
rent/finalized/subdir0/subdir0/blk_1073741850", O_RDONLY) = 319
31179 16:04:22 fstat(319, {st_mode=S_IFREG|0664, st_size=67108864, ...}) = 0

```

File content is sent by sendfile(), 64 KB at a time. At the same time, meta-data of 543 bytes are transferred through write() operation. Note that epoll\_wait() is called to wait the socket to be available for writing.

```

31179 16:04:22 epoll_ctl(328, EPOLL_CTL_ADD, 316, {EPOLLOUT, {u32=316,
u64=16965504754916524348}}) = 0
31179 16:04:22 epoll_wait(328, [{EPOLLOUT, {u32=316, u64=16965504754916524348}}], 8192,
480000) = 1
31179 16:04:22 epoll_ctl(328, EPOLL_CTL_DEL, 316, 0x7f965d090180) = 0
31179 16:04:22 epoll_wait(328, [], 8192, 0) = 0
31179 16:04:22 fstat(319, {st_mode=S_IFREG|0664, st_size=67108864, ...}) = 0
31179 16:04:22 sendfile(316, 319, [0] => [65536], 65536) = 65536
31179 16:04:22 write(316, "\0\1\2\4\0\31\t\0\0\1\0\0\0\0\0\2\1\1\0\0\0\0\0\0\30\0%\0\0\1\0\0"..., 543)
= 543
...

```

```
#!/bin/bash
for num in {1..100}
do
    echo 'hello world' | hdfs dfs -put - /test/$num.txt
done
```

- 1) NameNode:
  - a) Most frequently used system calls

According to the above statistic, The most frequently used system calls related to file system are 2697 read(), 1728 write(), 527 lseek(), 400 fdatasync().

```
4621 22:48:56 lseek(208, 0, SEEK_CUR) = 101534
4621 22:48:56 fstat(208, {st_mode=S_IFREG|0664, st_size=1048576, ...}) = 0
4621 22:48:56 write(208,
"\0\0\0\0\0253\0\0\0\0\0\0\04\323\0\0\0\0\0@ \323\0\25/test/2.txt._COPYING_\0\1\0\0\1[\27\367Uq\0\0\1[\27\367Uq\0\0\0\0\10\0\0\0\0\0\0\10holmesin\nsupergroup\1\244\0\0\0\0"... , 176) = 176
4621 22:48:56 fdatsync(208)          = 0
.....
4627 22:48:56 lseek(208, 0, SEEK_CUR) = 101710
4627 22:48:56 fstat(208, {st_mode=S_IFREG|0664, st_size=1048576, ...}) = 0
4627 22:48:56 write(208, "
```

```

\0\0\0\24\0\0\0\0\0\0\0\4\324\0\0\0\0\0@ \0\0\0\314\373S8\10\37\0\0\0\24\0\0\0\0\0\0\4\325\0\0\0\0\0\0\4\2
64\335\267\360f!\0\0\0006\0\0\0\0\0\0\0\4\326\0\25/test/2.txt._COPYING_\1\0\0\0\0@ \0\0\0\314\0\216\
4\264\0"... , 109) = 109
4627 22:48:56 fdatsync(208)          = 0
.....
4624 22:48:57 lseek(208, 0, SEEK_CUR) = 101819
4624 22:48:57 fstat(208, {st_mode=S_IFREG|0664, st_size=1048576, ...}) = 0
4624 22:48:57 write(208,
"\f\0\0\0w\0\0\0\0\0\0\0\4\327\0\0\0\0\0\0\0\0\0\25/test/2.txt._COPYING_\0\1\0\0\1[\27\367V\361\0\0\1
[\27\367Uq\0\0\0\0\10\0\0\0\0\0\0\1\0\0\0\0\0@ \0\0\0\314\0\0\0\0\0\0\0\0\0\0\0\0\4\264\10h"... , 124)
= 124
4624 22:48:57 fdatsync(208)          = 0
.....
4623 22:48:57 lseek(208, 0, SEEK_CUR) = 101943
4623 22:48:57 fstat(208, {st_mode=S_IFREG|0664, st_size=1048576, ...}) = 0
4623 22:48:57 write(208,
"\1\0\0\0N\0\0\0\0\0\0\0\4\330\0\25/test/2.txt._COPYING_\0w/test/2.txt\0\0\1[\27\367V\370\0\20U\177
\17\25\305A\251Y;\322*pQ\0\0\0\0\7f\275\2128", 83) = 83
4623 22:48:57 fdatsync(208)          = 0

```

After creation of the file on the datanode, the namenode will append the metadata to the file with the file handler 208, the file created has a suffix “\_COPYING\_”. This metadata is flushed to disk. Then after block added on the datanode, metadata will be appended to file 208, and flushed to disk. After the completion of the file creation, metadata will be flushed to disk for file 208. Then file will be renamed without the “\_COPYING\_” suffix, this metadata is write to file 208, and flushed to disk. To conclude, for each creation of a small file, it involves 4 metadata write and fdatsync() operation.

#### c) Finding:

After three times run of the 100 creation process, some different file access pattern was found.

```

4620 22:51:01 write(191, "2017-03-28 22:51:01,667 INFO
org.apache.hadoop.hdfs.server.namenode.FSEditLog: Starting log segment "... , 108) = 108
4620 22:51:01
open("/tmp/hadoop-holmesin/dfs/name/current/edits_inprogress_0000000000000001500",
O_RDWR|O_CREAT, 0666) = 208

```

During the operation, some operation relates to rolling edit logs happens. And the file above was created with file handle 208. After observation, all the new metadata change will try to write to this file. The log information in the above table implies the metadata on the namenode is log based. And each metadata change is written to the file “edits\_inprogress\_xxx”.

#### 2) DataNode:

##### a) Most frequently used system calls:

Strace with -c option result:

% time	seconds	usecs/call	calls	errors	syscall
48.13	0.196000	32667	6		poll
48.13	0.196000	1960	100		accept
2.76	0.011225	12	904		write
0.98	0.004000	10	414		mmap
0.00	0.000000	0	860	401	read
0.00	0.000000	0	304		open

0.00	0.000000	0	404	close
0.00	0.000000	0	309	107 stat
0.00	0.000000	0	304	fstat
0.00	0.000000	0	42	lseek
0.00	0.000000	0	2945	mprotect
0.00	0.000000	0	809	rt_sigprocmask
0.00	0.000000	0	551	sched_yield
0.00	0.000000	0	200	madvise
0.00	0.000000	0	100	dup2
0.00	0.000000	0	1	socket
0.00	0.000000	0	1	1 connect
0.00	0.000000	0	202	getsockname
0.00	0.000000	0	102	setsockopt
0.00	0.000000	0	1	getsockopt
0.00	0.000000	0	203	clone
0.00	0.000000	0	306	fcntl
0.00	0.000000	0	200	rename
0.00	0.000000	0	2	unlink
0.00	0.000000	0	203	statfs
0.00	0.000000	0	203	gettid
0.00	0.000000	0	406	sched_getaffinity
0.00	0.000000	0	804	epoll_ctl
0.00	0.000000	0	203	set_robust_list
-----				
100.00	0.407225		11089	509 total

According to the above statistic, The most frequently used system calls related to file system are 904 write(), 860 read(), 304 open(), 42 lseek(), 200 rename().

b) File system access pattern:

```

20732 22:48:57
stat("/tmp/hadoop-holmesin/dfs/data/current/BP-1056811197-127.0.1.1-1490748804937/current/rbw/blk_1073742028", 0x7ff10ebd3200) = -1 ENOENT (No such file or directory)
20732 22:48:57
open("/tmp/hadoop-holmesin/dfs/data/current/BP-1056811197-127.0.1.1-1490748804937/current/rbw/blk_1073742028", O_RDWR|O_CREAT|O_EXCL, 0666) = 240
20732 22:48:57 fstat(240, {st_mode=S_IFREG|0664, st_size=0, ...}) = 0
20732 22:48:57 close(240) = 0
20732 22:48:57
open("/tmp/hadoop-holmesin/dfs/data/current/BP-1056811197-127.0.1.1-1490748804937/current/rbw/blk_1073742028_1204.meta", O_RDWR|O_CREAT, 0666) = 240
20732 22:48:57 fstat(240, {st_mode=S_IFREG|0664, st_size=0, ...}) = 0
20732 22:48:57
open("/tmp/hadoop-holmesin/dfs/data/current/BP-1056811197-127.0.1.1-1490748804937/current/rbw/blk_1073742028", O_RDWR|O_CREAT, 0666) = 241
.....
20732 22:48:57 read(239,
"\0\0\0\24\0\31\t\0\0\0\0\0\0\0\21\0\0\0\0\0\0\0\30\0%\f\0\0\0\360\377r\222hello world\n", 512) =
47
20732 22:48:57 write(241, "hello world\n", 12) = 12
.....
20733 22:48:57
rename("/tmp/hadoop-holmesin/dfs/data/current/BP-1056811197-127.0.1.1-1490748804937/current/rbw/blk_1073742028_1204.meta",

```

```
"/tmp/hadoop-holmesin/dfs/data/current/BP-1056811197-127.0.1.1-1490748804937/current/finalize
d/substr0/substr0/blk_1073742028_1204.meta") = 0
20733 22:48:57
rename("/tmp/hadoop-holmesin/dfs/data/current/BP-1056811197-127.0.1.1-1490748804937/curren
t/rbw/blk_1073742028",
"/tmp/hadoop-holmesin/dfs/data/current/BP-1056811197-127.0.1.1-1490748804937/current/finalize
d/substr0/substr0/blk_1073742028") = 0
```

By observing the strace result on datanode, it will first create two file blk\_xxx and the metadata file for it blk\_xxx.meta under the folder .../rbw(Replica Being Written). Then it will read the "hello world" from socket, then write this string to the block file "blk\_xxx". After the write operation completed, the datanode will rename the two file from folder .../rbw to the finalized folder ".../finalized/...".

### 3.6 File system access pattern when creating folders:

```
#!/bin/bash
for num in {1..2}
do
    hdfs dfs -mkdir /test$num
done
```

The test try to create 100 new folders with the name like /testxx

- 1) NameNode:
  - a) Most frequently used system calls:  
from strace with -c result, it has 1402 read(), 154 lseek(), 517 write(), and 100 fdatasync().
  - b) Access pattern: for one mkdir operation, the metadata is written to "edits\_inprogress\_xxx", and flushed to disk.
- 2) DataNode:
  - a) Most frequently used system calls: from strace with -c result, there are 179 newfstatat(), 12 open(), 7 read(), 3 write().
  - b) Access pattern: When creating folders doesn't have much file system calls.

### File system access pattern when change metadata:

```
#!/bin/bash
for num in {1..100}
do
    hdfs dfs chmod 777 /test/$num.txt
done
```

The test try to change 100 files with a different access permission

- 3) NameNode:
  - a) Most frequently used system calls:  
from strace with -c result, it has 1145 read(), 415 write(), 122 lseek(), and 100 fdatasync().
  - b) Access pattern: for one chmod operation, namenode write() the change to the socket, then the new metadata is appended to "edits\_inprogress\_xxx", and flushed to disk. Notice the lseek() before fdatasync() try to locate the file handle to the current position, each metadata change in appended to the end of the "edits\_inprogress\_xxx" file, which confirms that the metadata is not overwritten, and it is log based.
- 4) DataNode: Changing of file access permission doesn't involve datanode operations in file system calls.



Access pattern:

[illegible]

```

4604 22:51:02 close(222)          = 0
.....
4604 22:51:02
open("/tmp/hadoop-holmesin/dfs/name/current/fsimage_0000000000000001499.md5.tmp",
O_WRONLY|O_CREAT|O_TRUNC, 0666) = 222
.....
4604 22:51:02
rename("/tmp/hadoop-holmesin/dfs/name/current/fsimage_0000000000000001499.md5.tmp",
"/tmp/hadoop-holmesin/dfs/name/current/fsimage_0000000000000001499.md5") = 0
4604 22:51:02
rename("/tmp/hadoop-holmesin/dfs/name/current/fsimage.ckpt_0000000000000001499",
"/tmp/hadoop-holmesin/dfs/name/current/fsimage_0000000000000001499") = 0
.....

4604 22:51:02 unlink("/tmp/hadoop-holmesin/dfs/name/current/fsimage_000000000000000000")
= 0
4604 22:51:02
unlink("/tmp/hadoop-holmesin/dfs/name/current/fsimage_000000000000000000.md5") = 0

```

When the rolling of edit log happens, the old “edits\_inprogress\_000000000000000006” is renamed to a edits log file with start and end number “edits\_000000000000000006-0000000000000001499”. And a new “edits\_inprogress\_xxx” file is created. New “fsimage.ckpt\_xxx”(ckpt stands for checkpoint) and “fsimage\_xxx.md5.tmp” are created. Then the namenode try to read from the old edits log file, and then write those log to a new “fsimage\_xxx” file. After merging the old fsimage and edits log file, The old fsimage file is then deleted. Fimage file and md5 file will be renamed.

Checkpointing is the process of merging the most recent fsimage with all edits applied after that fsimage is merged in order to create a new fsimage. After looking into the file hdfs-default.xml

```

<property>
  <name>dfs.namenode.checkpoint.period</name>
  <value>3600</value>
  <description>The number of seconds between two periodic checkpoints.
</description>
</property>

<property>
  <name>dfs.namenode.checkpoint.txns</name>
  <value>1000000</value>
  <description>The Secondary NameNode or CheckpointNode will create a checkpoint
of the namespace every 'dfs.namenode.checkpoint.txns' transactions, regardless
of whether 'dfs.namenode.checkpoint.period' has expired.
</description>
</property>

```

The checkpointing process is triggered if time has elapsed 3600 since last checkpoint, or if current transactions in edits file has exceeded 1000000.

### 3.8. Possible optimization:

By observation, the namenode will store all the metadata transaction to the file onto the disk. Even for creating a small file, it invokes four fdatasync() I/O. These disk I/O can take a lot of time, which affects

the response time to client. One possible optimization would be, batch some writes within a single disk I/O, since these metadata log are continuous in the file.